

Feature Vectors

Content-based Image Retrieval

Image Features

- Low-level feature extraction
 - How to represent an image in a compact and descriptive way?
 - How to compare features, and, thus, images?

Image Features

Image Database



Images





Image Feature
Extraction



Feature Space

Query by Visual Example

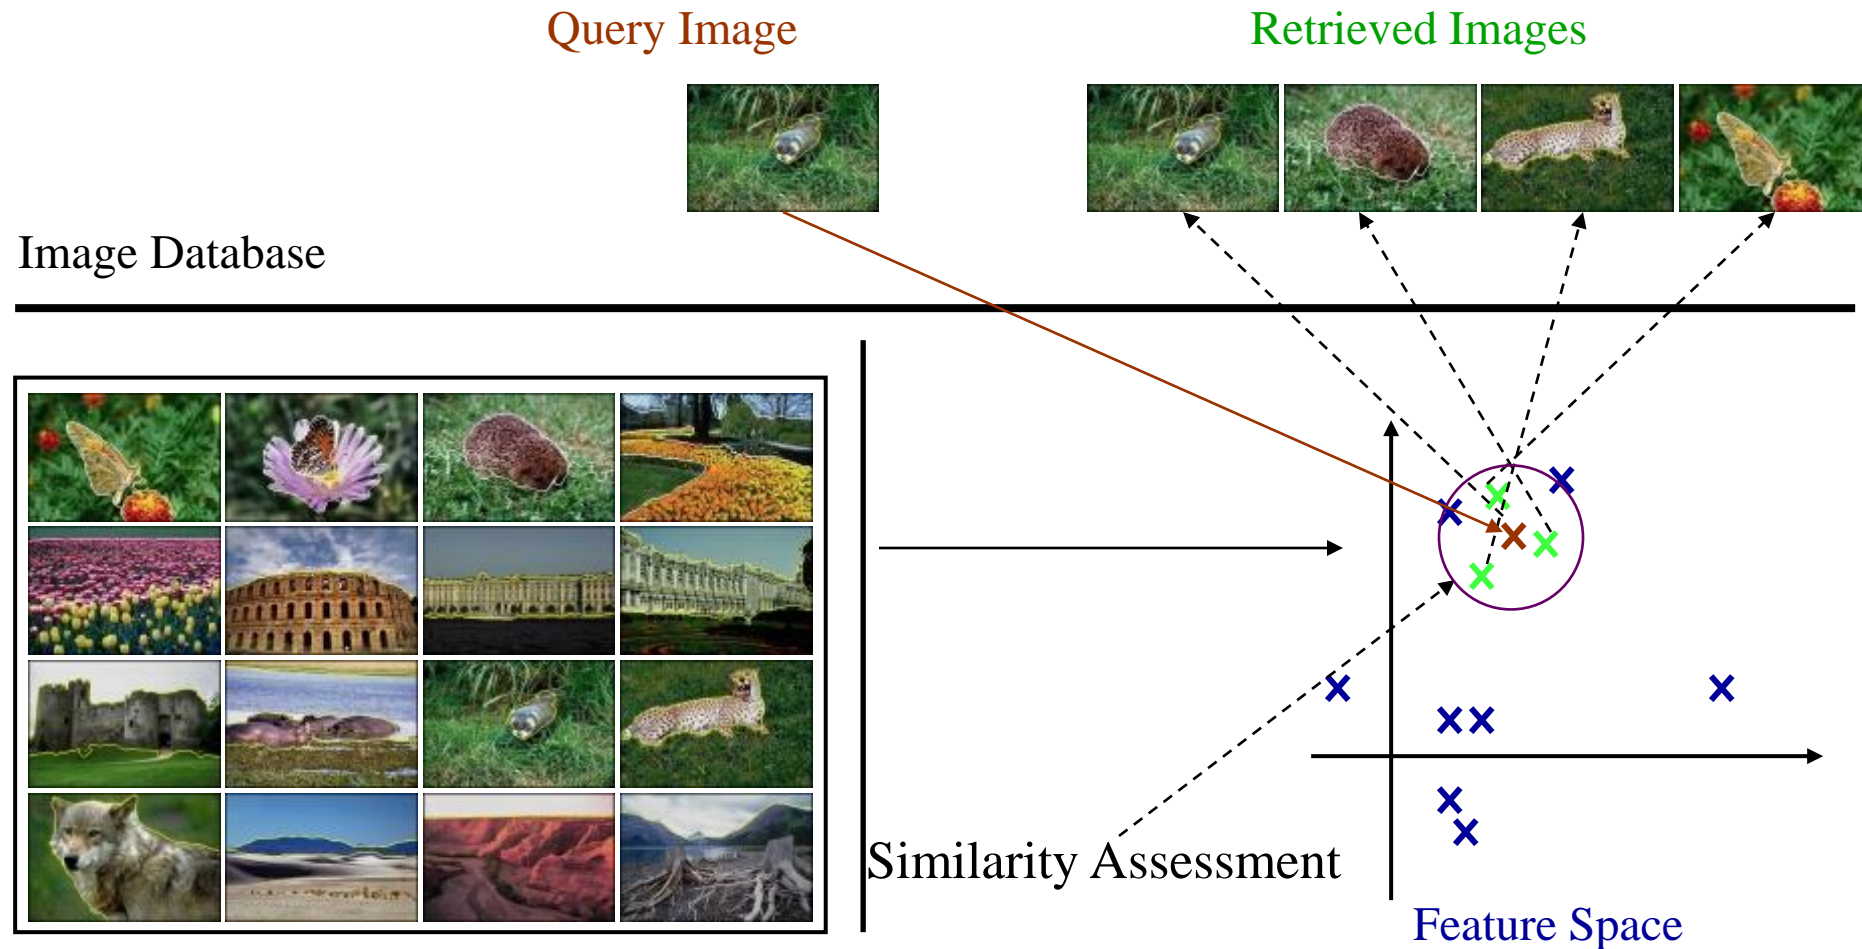


Image Features

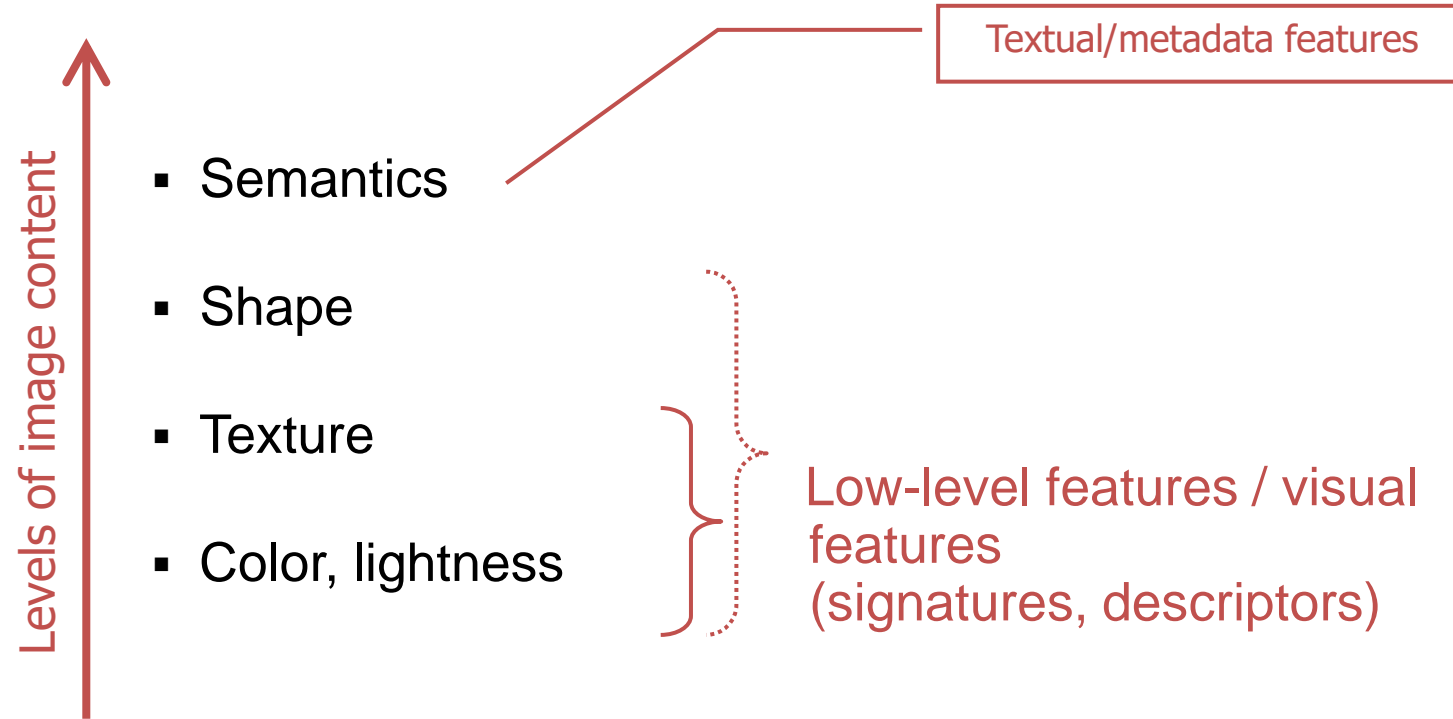



Image Features




Annotations and metadata:

- tags/keywords;
- creation date;
- geo tags;
- name of the file;
- photography conditions (exposition, aperture, flash...).

Features extracted from pixel values:

- color descriptors;
- texture descriptors;
- shape descriptors;
- spatial layout descriptors.

Image Features



Describes **the whole image**:

- average intensity;
- average amount of red;
- ...

All pixels of the image are processed.

Describes **one part** of the image:

- average intensity for the left upper part;
- average amount of red in the center of the image;
- ...

Segmentation of the image is performed, pixels of a particular segment are processed to extract features, Key point detection and localization

Feature Spaces

- **Feature vector** – a vector of features, representing one image.
- **Feature space** – the set of all possible feature vectors with defined similarity measure.

Image A



x^A_1	x^A_2	...	x^A_N
---------	---------	-----	---------

y^A_1	y^A_2	...	y^A_M
---------	---------	-----	---------

z^A_1	z^A_2	...	z^A_K
---------	---------	-----	---------

⋮

Image B



x^B_1	x^B_2	...	x^B_N
---------	---------	-----	---------

y^B_1	y^B_2	...	y^B_M
---------	---------	-----	---------

z^B_1	z^B_2	...	z^B_K
---------	---------	-----	---------

⋮

Feature Spaces

- **Feature vector** – a vector of features, representing one image.
- **Feature space** – the set of all possible feature vectors with defined similarity measure.

Image A



x_1^A x_2^A ... x_N^A

y_1^A y_2^A ... y_M^A

z_1^A z_2^A ... z_K^A

⋮

Image B



x_1^B x_2^B ... x_N^B

y_1^B y_2^B ... y_M^B

z_1^B z_2^B ... z_K^B

⋮

← Similarity measure →

← Similarity measure →

← Similarity measure →


Feature Spaces

- **Feature vector** – a vector of features, representing one image.
- **Feature space** – the set of all possible feature vectors with defined similarity measure.

Image A



Image B

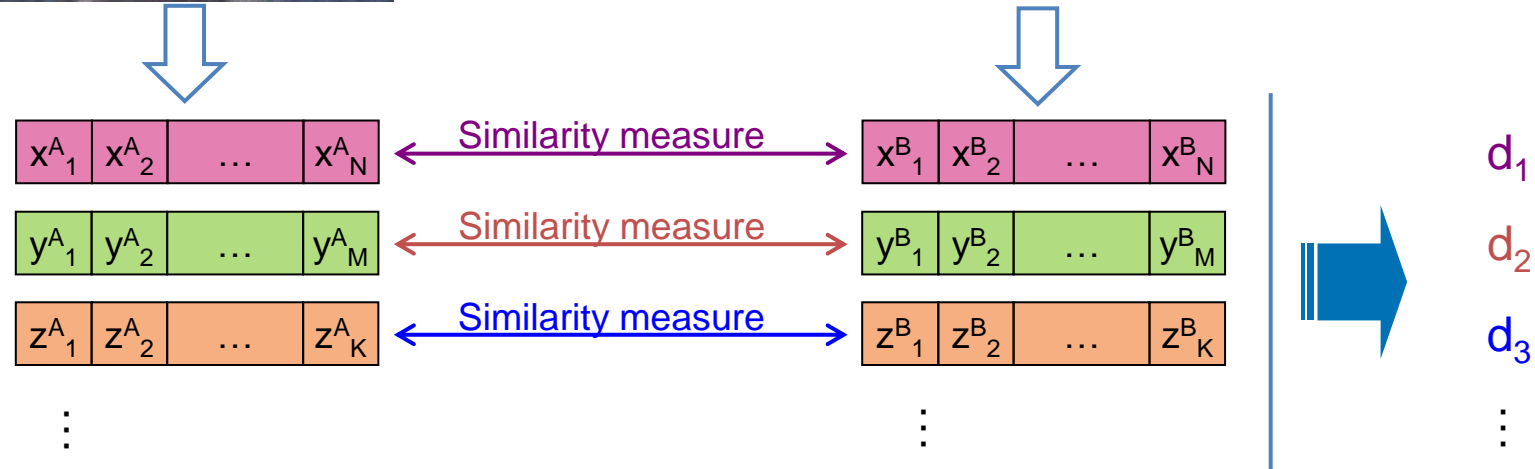


Combine Results

Image A



Image B



Content-based Image Retrieval

- General: We try to match a query image to an arbitrary collection of images, such as those found on the web. The goal of the query is to obtain images with the same object as the query
 - Given an image with a horse, find all images showing a horse (at least as their main subject)
- Application Specific: We try to match a query image to a collection of images of a specific type.
 - For example, fingerprints, X-ray images of a specific organ, etc.

Content-based Image Retrieval

- Query: image/video clip Retrieve: images/shots from archive



near duplicate



same object



same category

Color Histograms for CBIR

- Histograms can serve as feature vectors (i.e. a list of numbers used to quantify an image and compare it to other images)
- Assumption: images with similar color distributions are semantically similar
- Comparing the “similarity” of color histograms can be done using a distance metric
 - Euclidean, correlation, Chi-squared, intersection, and Bhattacharyya

Using OpenCV to Compute Histograms

- `cv2.calcHist(images, channels, mask, histSize, ranges)`
- **images:** This is the image that we want to compute a histogram for. Wrap it as a list: `[myImage]`
- **channels:** A list of indexes, where we specify the index of the channel we want to compute a histogram for. To compute a histogram of a grayscale image, the list would be `[0]`. To compute a histogram for all three red, green, and blue channels, the channels list would be `[0, 1, 2]`
- **mask:** `uint8` image with the same shape as our original image, where pixels with a value of zero are ignored and pixels with a value greater than zero are included in the histogram computation. Using masks allow us to only compute a histogram for a particular region of an image
- **histSize:** This is the number of bins we want to use when computing a histogram. Again, this is a list, one for each channel we are computing a histogram for. The bin sizes do not all have to be the same. Here is an example of 32 bins for each channel: `[32, 32, 32]`
- **ranges:** The range of possible pixel values. Normally, this is `[0, 256]` for each channel, but if you are using a color space other than RGB (such as HSV), the ranges might be different

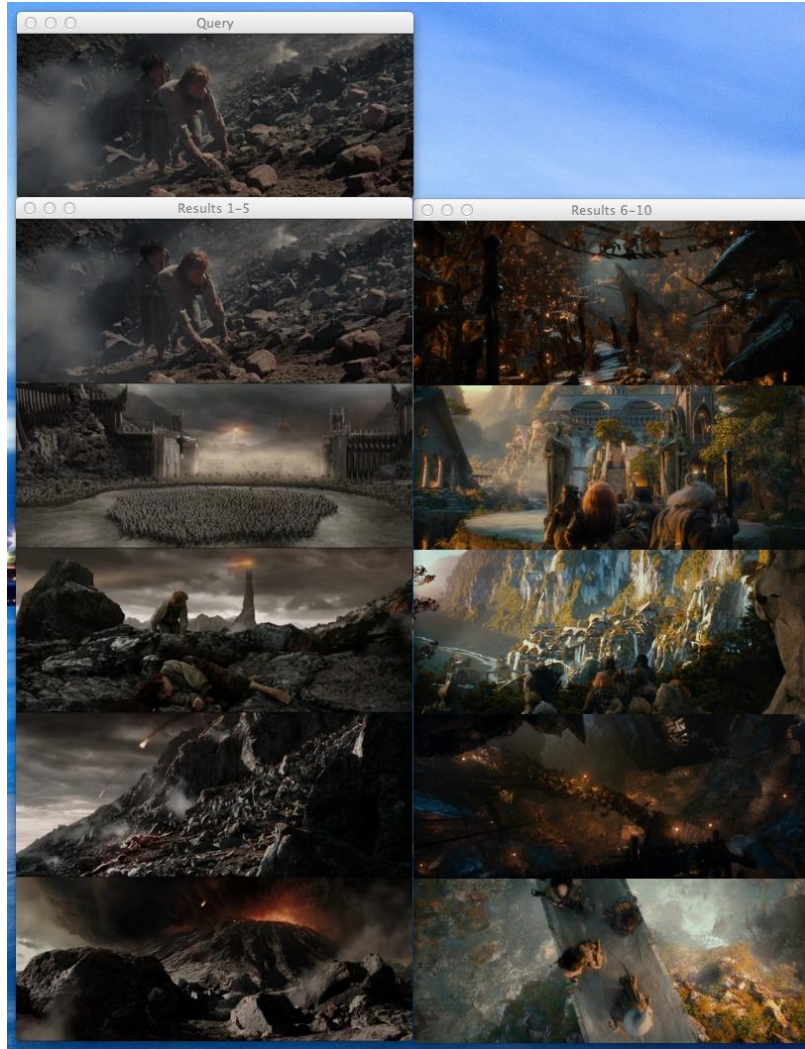
Compare Histograms

- `cv2.compareHist(H1, H2, method)`
 - `cv2.cv.CV_COMP_CORREL`: Computes the correlation between the two histograms
 - `cv2.cv.CV_COMP_CHISQR`: Applies the Chi-Squared distance to the histograms

$$\frac{1}{2} \sum_{i=1}^n \frac{(x_i - y_i)^2}{(x_i + y_i)}$$

- `cv2.cv.CV_COMP_INTERSECT`: Calculates the intersection between two histograms
- `cv2.cv.CV_COMP_BHATTACHARYYA`: Bhattacharyya distance, used to measure the “overlap” between the two histograms
- `cv2.cv.CV_COMP_HELLINGER`: A synonym for `cv2.cv.CV_COMP_BHATTACHARYYA`

Example: Query by Color



Problems

- Two Images with nearly identical color histograms



Problems

- Perceptual versus computational similarity

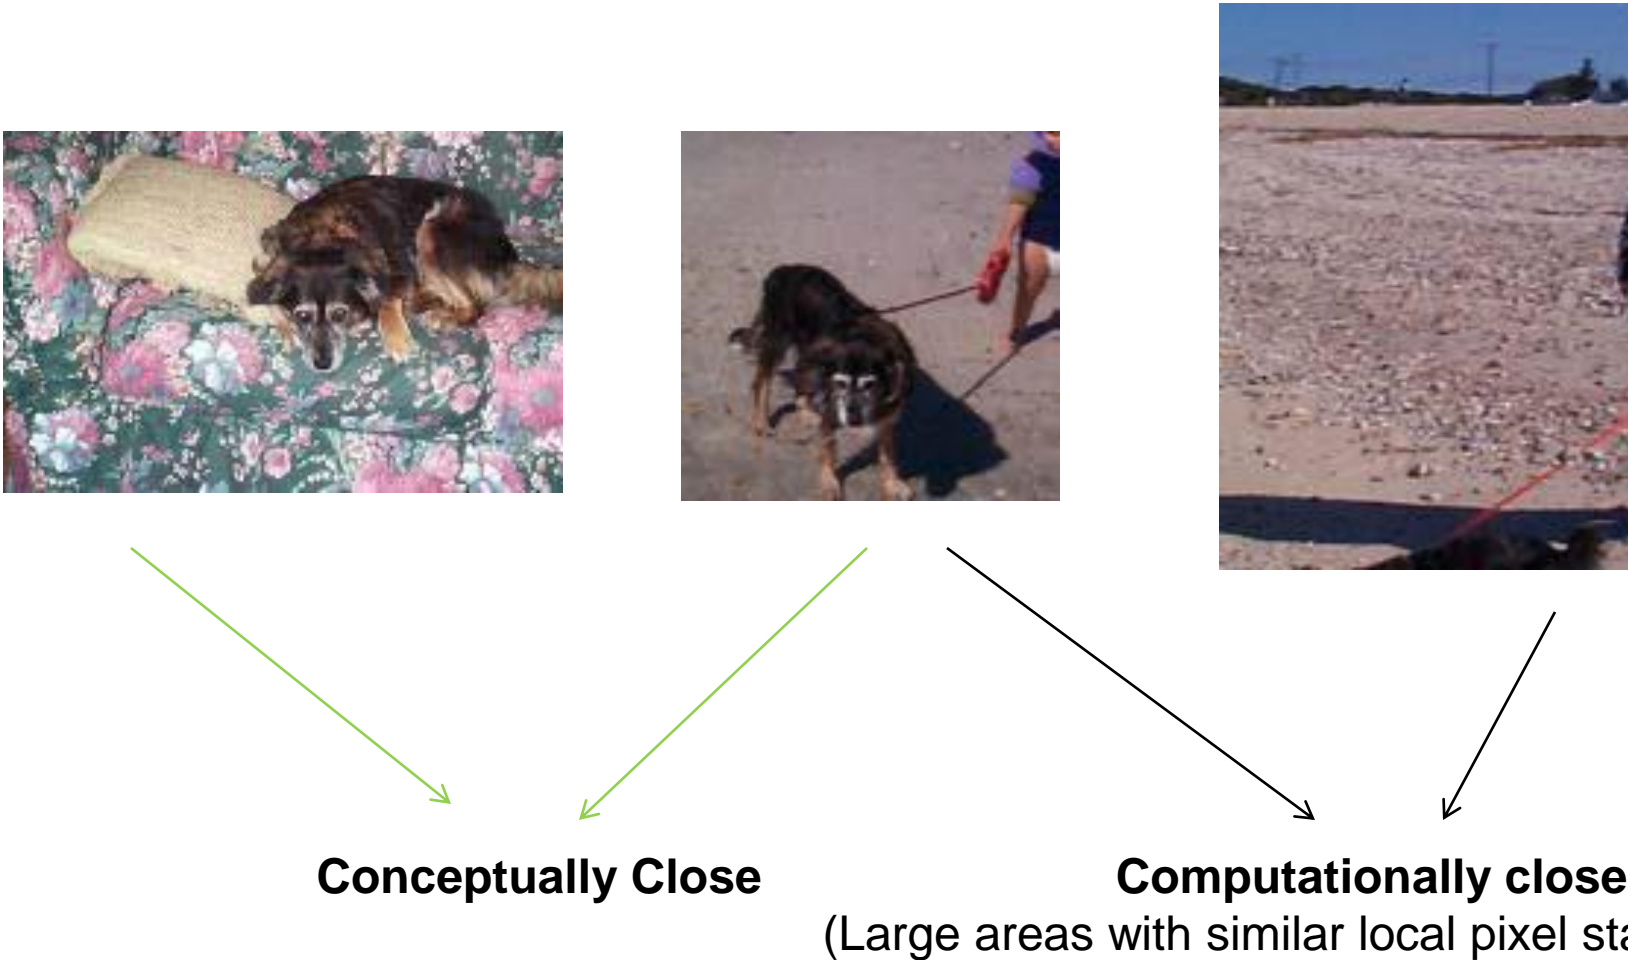


Perceptually close

**Computationally close
(similar pixel statistics)**

Problems

- Conceptual similarity is even harder to deal with



Problems


- Particular objects, not entire images



- Forced to face problems of:
 - scale change
 - pose change
 - illumination change
 - partial occlusion

Possible solution

- When do (images of) objects match?



- Two requirements:
 - “patches” (parts) correspond, and
 - Configuration (spatial layout) corresponds